

Is Code Better Than Language for Algorithmic Reasoning?

Terry Tong¹ Yu Feng¹ Surbhi Goel¹ Dan Roth¹

Abstract

For tool-augmented language models, direct comparison between natural-language reasoning and code-execution pipelines are difficult to interpret because they conflate the intermediate representation and execution mechanisms. We resolve this ambiguity by inserting an intermediate intervention: the model expresses its reasoning as executable code, which is simulated in-context by the language model to produce an answer. On a 40-task verifiable algorithmic analysis set, deterministic code execution outperforms natural-language reasoning by +31.6pp, while in-model code simulation is not meaningfully different from natural-language reasoning (+0.15pp; 95% cluster-bootstrap CI $[-0.30, +0.61]$ pp). This suggests that, in our evaluated setting, changing the intermediate representation alone does not sufficiently explain the tool-use advantage, and that performance gains instead arise in conjunction with reliable external execution. Supporting this interpretation, a reconstruction intervention experiment that uses a proxy language model to infer natural-language reasoning traces from code representations recovers performance comparable to the original natural-language reasoning pipeline. We formalize this intuition with a simple statistical decision-theoretic model that characterizes when execution dominates the end-to-end risk in our disentangled trace generation-execution regime. All experiments are at <https://github.com/TerryTong-Git/ToolProj>.

1. Introduction

Many agentic systems orchestrate symbolic solvers, LLMs, and other tools, demonstrating state-of-the-art performance (Gao et al., 2023; Yao et al., 2023; Schick et al., 2023; Yang et al., 2024; Wang et al., 2025). Prior research em-

^{*}Equal contribution ¹University of Pennsylvania. Correspondence to: Terry Tong <tongt1@seas.upenn.edu>.




Figure 1. Three-Route Framework. We decompose algorithmic reasoning into: (1) *Translation* into NL or code, and (2) *Execution* via LLM or solver. This yields three routes: **Route 1** (Direct NL), **Route 2** (Code + NL simulation), **Route 3** (Code + Solver Execution). Prior work compares only Route 1 vs. Route 3, confounding translation and execution. Our Route 2 isolates these factors.

pirically demonstrates that translating problems into solver-executable code representations (hereafter denoted Route 3) and delegating execution often outperforms reasoning end-to-end in NL (hereafter denoted Route 1) on logic- and algorithmic-style complex reasoning tasks (Lyu et al., 2023; Pan et al., 2023). However, it is unclear whether benefits come from the more structured code representation itself used in the planning and reasoning process, or the reliability of external solvers, or both. Little progress has been made in this direction since the problem itself is ill-posed: two routes learn fundamentally different objects, so there is no common formal target and metric to compare.

This paper presents a systematic three-route framework (Figure 1 and Section 2) that breaks down the question into tractable pairwise comparisons between the different routes. We decouple the end-to-end pipeline into a trace generation phase (e.g. Chain of Thought (Wei et al., 2022)) and an execution phase (see Section 2.2), instantiated by controlled prompts in Figure 2. This allows us to fix one phase, and make causal claims on the effects of phase on end-to-end reasoning performance. We first instantiate this framework and find the empirical pattern Route 1 (NL + NL Reasoning) $\approx$ Route 2 (Code + NL Reasoning) < Route 3 (Code + Python Execution) (Sections 3 and 3.2). On the 40-task analysis set, Route 2 is +0.15pp above Route 1 with a 95% cluster-bootstrap CI crossing zero, while Route 3 exceeds Route 1 by +31.6pp and Route 2 by +31.5pp across six models and seeds $\{0, 1, 2\}$. Our results suggest Route 3 (Code + Python Execution) is optimal (Figures 3 and 4). To under-

Route 1 — NL	Route 2 — Code + sim.	Route 3 — Code + Python
PROMPT Solve the algorithmic problem step by step in natural language. Return JSON: {simulation, answer}.	PROMPT Solve the algorithmic problem. Return code + a NL simulation. Return JSON: {code, simulation, answer}.	PROCEDURE 1. Extract solution() from Route 2. 2. Run in sandboxed Python (5 s). 3. Compare output to gold.
CONSTRAINT no code	CONSTRAINT same JSON schema	NOTE no additional LLM call
EXAMPLE OUTPUT · Q: 123 + 456 <pre>{ "simulation": "3+6=9, 2+5=7, 1+4=5", "answer": "579" }</pre>	EXAMPLE OUTPUT · Q: 123 + 456 <pre>{ "code": "def solution(): return 123+456", "simulation": "adds", "answer": "579" }</pre>	EXECUTION · CODE FROM ROUTE 2 <pre>>>> def solution(): ... return 123 + 456 >>> solution() 579</pre>
answer 579	answer 579	answer 579

Figure 2. Prompt templates for three-route evaluation. **Route 1 (NL)**: LLM reasons in natural language only, code forbidden. **Route 2 (Sim)**: LLM generates Python `solution()` then simulates execution in NL. **Route 3 (Code)**: Same code executed in Python runtime. This isolates execution mechanism while controlling translation.

stand where language is bottlenecked, we operationalize our framework below.

When analyzing the pair Route 1 v.s. Route 2 (Section 3.3), we fix the execution phase to be a LLM with NL reasoning, and vary the representation modality of its Chain-of-Thought to be either code or NL. We first model the linear case, showing that under the assumptions that code is more structured than natural language in the nuisance parameter sense, then the worst case risk for code representations are uniformly better than natural language representations over all hypotheses in our linear hypothesis class (Section 3.3.1 and Theorem 3.5). We show that this fixed transformation nuisance perspective holds in the non-linear and high-dimensional LLM regime when the theory breaks down (Section 3.4 and Figure 5). The results show that code reasoning traces, when transformed to natural language reasoning via a proxy LM, do just as well as the original natural language traces. This provides evidence against representation as the bottleneck.

We further analyze Route 2 v.s. Route 3 (Section 4), identifying execution to be the bottleneck. In the 40-task analysis set, the recovery mass where Route 2 succeeds while Route 3 fails is 1.61%, while the execution-win mass where Route 3 succeeds while Route 2 fails is 33.08% (Figure 6). On an analysis of hardness scaling, we find that code execution (Route 3) scales much better as problems become more difficult.

Our main contributions are:

1. A **three-route framework** (Section 2 and Figures 1 and 2) for tractable comparison between code and natural language representations via an intermediary (code generation with LLM execution).
2. **Empirical validation** (Section 3 and Figures 3 and 4) demonstrating that Route 1 and Route 2 are statistically close, while Route 3 is substantially better.
3. A **systematic study** (Sections 3.3, 3.4 and 4) ruling out trace generation, and solidifying execution as the bottleneck in end-to-end algorithmic reasoning performance when using language.

2. Three-Route Framework

Our central research question (RQ) is: *Is code > NL for algorithmic reasoning?* To begin to answer this question, we first introduce our three-route framework which enables tractable comparison by disentangling *reasoning representation* (hereafter called *Traces*) and *reasoning execution* (hereafter called *Executors*) and constructing an intermediary bridge (Route 2) for pairwise comparison. This section introduces the task (Section 2.1), notation (Section 2.2), and route definitions (Section 2.3).

2.1. Task and Loss

For a gold evaluation distribution over task instances i specified by (algorithm, input variables, seed), let $X \sim p(x)$ denote a task instance (problem statement and inputs), and let $Y^*(X) \in \mathcal{Y}$ denote the ground-truth output that is unique, fixed, and externally verifiable. We evaluate performance under 0–1 loss:

$$\ell(y, x) := \mathbf{1}\{y \neq Y^*(x)\}.$$

2.2. Trace Generators and Executors

We model each reasoning pipeline as a two-stage stochastic process with (1) *trace generation* and (2) *execution*.

Trace generation. We define a *trace* as an object that stores intermediate reasoning used to solve a corresponding task (e.g. NL Chain-of-Thought or a program). A *trace generator* is a (stochastic) mapping

$$E : \mathcal{X} \rightarrow \Delta(\mathcal{Z}), \quad Z \sim p_E(z | x),$$

which produces an auxiliary trace Z given the task instance.

Execution. We define *execution* as a procedure that consumes a trace, e.g. an LLM forward pass that takes as input a reasoning trace and outputs an answer, or executing a program in an external runtime. An *executor* is a (stochastic) mapping

$$\rho : \mathcal{X} \times \mathcal{Z} \rightarrow \Delta(\mathcal{Y}),$$

mapping the observed instance and trace to a final output.




Figure 3. **Code + Solver Execution performs better than Direct NL reasoning quantified by end-task accuracy overall and within paired instances.** Paired instance contrasts are shown in the bootstrap distributions. Results are averaged over the 40-task analysis set with 1,113 unique problem instances, 3 seeds, and 6 models.

2.3. The Three Routes

We consider three reasoning pipelines (“routes”), illustrated in Figure 1. Each route is represented as a pair (E, ρ) consisting of a trace generator E and an executor ρ .

Route 1 (Direct Natural Language). Route 1 represents standard NL reasoning: the model first produces a natural-language (Chain-of-thought) trace and then the same model is used as the LLM-based executor, conditioning on the trace to generate a final answer. Formally, a natural-language trace generator E_{NL} produces traces $Z_{NL} \sim p_{NL}(\cdot | X)$, paired with an executor family

$$\mathcal{H}_{NL} \subseteq \{\rho : \mathcal{X} \times \mathcal{Z}_{NL} \rightarrow \Delta(\mathcal{Y})\}.$$

Route 2 (Code + NL Simulation). Route 2 uses *code* as the trace modality, but keeps execution “in-model”: the LLM instead simulates the generated code in natural language, rather than executing it in an external environment. This route isolates the effect of using a code trace while holding the LLM-based executor class fixed. Formally, a code trace generator E_{Code} produces executable representations $Z_C \sim p_{Code}(\cdot | X)$, paired with an executor family

$$\mathcal{H}_C \subseteq \{\rho : \mathcal{X} \times \mathcal{Z}_C \rightarrow \Delta(\mathcal{Y})\}$$

corresponding to language-model-based simulation of code execution.

Route 3 (Code + Solver Execution). Route 3 uses the same code trace generator E_{Code} as Route 2, producing $Z_C \sim p_{Code}(\cdot | X)$, but pairs it with a deterministic executor corresponding to external code execution. Let $\text{Exec} : \mathcal{X} \times \mathcal{Z}_C \rightarrow \mathcal{Y}$ denote the (fixed) external runtime that executes the code trace z on instance x and returns an output. Our executor is

$$\rho_{\text{Exec}}(y | x, z) := \mathbf{1}\{y = \text{Exec}(x, z)\}.$$

The corresponding executor family is the singleton $\mathcal{H}_{\text{Exec}} := \{\rho_{\text{Exec}}\}$.

3. Evaluating the Three-Route Framework

Our evaluation aims to answer the research questions (RQ):
1) Does Route 1 $\approx$ Route 2 < Route 3 hold? (§3.1, §3.2)

3.1. Experimental Instantiation of the Three Routes

To draw conclusions about the effect of modality in the trace generation or different execution methods, we ensure one stage is always fixed. For Route 1 vs Route 2, we fix the execution phase by using the same LLM reasoning model forward pass, but use different modalities in the trace generator. For Route 2 vs Route 3, we fix the trace generator which outputs code, then use different execution methods, either LLM reasoning model forward pass, or a python3 runtime. Below, let X_i be the task instantiation of instance i . Structured JSON outputs are enforced for sampling in all routes (Figure 2).

Sampling Route 1. Route 1 prompts an LLM E_{NL} to function as a trace generator $Z_{NL} := E_{NL}(X_i)$. The trace is then fed into the same LLM $\rho_{NL} = E_{NL}$ to produce an answer $Y_i^{(NL)} := \rho_{NL}(Z_{NL})$. The prompt instructs the model to never use code, and to output a structured rationale and answer (see Figure 2). Operationalized, one experimental run for Route 1 is encapsulated in a single LLM reasoning model’s forward pass without pause.

Sampling Route 2. Similarly, Route 2 uses a LLM E_{Code} as a trace generator for code modality $Z_{Code} := E_{Code}(X_i)$. The trace is then fed into the same LLM $\rho_{Sim} = E_{Code}$ to produce an answer $Y_i^{(Sim)} := \rho_{Sim}(Z_{Code})$. Operationalized, one experimental run for Route 2 is encapsulated in a single LLM reasoning model’s forward pass without pause. Prompt templates are controlled to be as similar as possible, with Route 2 differing only by the inclusion of a code-generation and simulation instruction (Figure 2). We prompt the model to output a program snippet, structured reasoning over the code, followed by an answer.




Figure 4. **Code + Solver Execution scales better than natural language reasoning when problems get harder, both within-task and on average, interpolating and extrapolating (grey)**. τ is used to control digit length for arithmetic, table dimensionality for dynamic programming, and constraint matrix dimensionality for integer linear programming. The same data and setup is used from Figure 3.

Sampling Route 3. Route 3 takes the exact same code trace Z_{Code} as Route 2, except it runs it through a python3 runtime ρ_{Exec} , and retrieves the solution $Y_i^{(Exec)} = \rho_{Exec}(Z_{Code})$. The python3 runtime has access to 5 standard scientific libraries: *Scipy*, *Numpy*, *Pandas*, *PuLP*, and *Pytorch*.

Observed outcomes. For each problem instance i , we observe paired Bernoulli outcomes

$$(Y_i^{(NL)}, Y_i^{(Sim)}, Y_i^{(Exec)}).$$

Data and models. We evaluate a 40-task analysis set drawn from CLRS-30, NP-Hard-Eval, and a custom fine-grained evaluation suite, across three random seeds $\{0, 1, 2\}$. The set contains 1,113 unique problem instances and 20,034 paired route evaluations after pooling six full-coverage models. Tasks span arithmetic, dynamic programming, graph algorithms, string algorithms, geometry, sorting, and NP-hard optimization, with difficulty controlled by a parameter τ when applicable. We evaluate closed-source models (Claude Haiku 4.5, GPT-4o-mini, Gemini 2.0 Flash, Gemini 2.5 Flash) and open-source models (Mixtral-8x22b-Instruct, Codestral-2508).

3.2. End-to-End Performance Comparison

Pairwise statistical tests. We evaluate pairwise route differences using McNemar tests on paired Bernoulli outcomes. For Route 1 vs. Route 2, the null hypothesis is

$$\begin{aligned} H_0 : & \Pr(Y^{(NL)} = 1, Y^{(Sim)} = 0) \\ &= \Pr(Y^{(NL)} = 0, Y^{(Sim)} = 1), \end{aligned}$$

with an analogous null for Route 2 vs. Route 3. Effect sizes are reported as paired accuracy differences, e.g.,

$$\Delta_{Sim-NL} = \text{Acc}(\text{Sim}) - \text{Acc}(\text{NL}),$$

with 95% confidence intervals estimated via cluster bootstrap resampling over instances. Holm-Bonferroni correction is applied to control family-wise error rate of the above 2 marginal pairwise tests on fully pooled data at $\alpha = 5\%$.

Difficulty scaling. To analyze how performance varies with task difficulty, we fit a generalized linear mixed-effects model (GLMM) for binary accuracy $Y_i \in \{0, 1\}$ with a logistic link:

$$Y_i \mid u_{inst[i]}, u_{seed[i]} \sim \text{Bernoulli}(p_i),$$

$\text{logit}(p_i) = \alpha + \beta_{route_i} + \gamma \tau_i + \delta_{route_i} \tau_i + u_{inst[i]} + u_{seed[i]}$, where $u_{inst[i]} \sim \mathcal{N}(0, \sigma_{inst}^2)$ and $u_{seed[i]} \sim \mathcal{N}(0, \sigma_{seed}^2)$ are random intercepts. Route and difficulty are modeled as fixed effects (including a route $\times$ difficulty interaction), with instance and seed modeled as random effects.

Results. Across the 40-task analysis set (Figure 3), Route 1 reaches 17.21% accuracy, Route 2 reaches 17.37%, and Route 3 reaches 48.84%. *Route 3 has a statistically significant advantage over Route 2*, with a +31.47pp paired accuracy gap and a 95% cluster-bootstrap interval of $[+29.20, +33.71]$ pp. For Route 1 vs. Route 2, the paired gap is +0.15pp with a 95% cluster-bootstrap interval of $[-0.30, +0.61]$ pp and a two-sided McNemar $p = 0.39$, so we do not conclude a meaningful difference between natural-language reasoning and code simulation under this evaluation. Performance gaps widen with increasing task difficulty (Figure 4): Route 3 remains substantially higher as the language-based routes degrade.

3.3. Route 1 vs Route 2 Analysis

In Section 3.2, Route 1 and Route 2 achieved similar paired accuracies under our evaluated models and prompts. The key research question in this section is, thus: **Can we provide evidence that input representation is not the bottleneck to end-task performance?** We first model the theory




Figure 5. Translated NL has the same functionality as the original native NL quantified by end-task accuracy. When concatenating the reasoning traces to the prompt, we get much higher accuracy as opposed to the baseline with just the prompt, indicating the model is using the reasoning. There are statistically significant differences between the baseline and concatenated prompts, but no statistically significant difference between the Translated NL and Native NL concatenations. Source data is randomly sampled from existing data used to generate Figure 3, we use 1000 samples per model.

in the simple linear case, then show that even when theory breaks down in the high-dimensional non-linear case, the results still hold.

The main result is that when we model code distribution as a canonicalization of natural language distribution, we can show for our simple linear hypothesis class that the uniform worst-case risk over *all* hypotheses is exactly 0. In other words, under these conditions, “code” reasoning is non-inferior to “natural language” reasoning.

3.3.1. LINEAR MODELLING SETUP AND INTUITION

Controlling the downstream hypothesis class enables us to compare the effect of the two input representations. For convenience, the downstream hypothesis class is modelled as realizable and linear.

$$\mathcal{H}_{\text{lin}} = \{h_{a,v}(b, u) = a^\top b + v^\top u : a \in \mathbb{R}^{d_B}, v \in \mathbb{R}^{d_U}\}. \quad (1)$$

and only the input distribution changes P_C and P_N , both of which are part of the same feature space. Let $r \in \{C, N\}$, where C denotes code and N denotes native natural language. For each task, route r produces a random input:

$$S_r = (B, U_r) \sim P_r, \quad S_r \in \mathbb{R}^{d_B + d_U}. \quad (2)$$

The label satisfies

$$Y = \theta^\top B + \varepsilon, \quad \mathbb{E}[\varepsilon \mid B, U_r] = 0, \quad (3)$$

and the nuisance is conditionally centered,

$$\mathbb{E}[U_r \mid B] = 0. \quad (4)$$

The vector $B \in \mathbb{R}^{d_B}$ is the answer-relevant core, or the sufficient statistic (Lehmann & Casella, 1998). The vector $U_r \in \mathbb{R}^{d_U}$ contains route-specific surface variation. Prior research in NLP / LLMs have shown that models are sensitive to syntactic heuristics, prompting, lexical choice, etc (Gururangan et al., 2018; McCoy et al., 2019; Geirhos et al., 2020; Zhao et al., 2021; Lu et al., 2022; Pezeshkpour & Hruschka, 2023).

Assumption 3.1 (Extra native-language nuisance). Conditional on the same core B , native natural-language nuisance decomposes as

$$U_N = U_C + \eta_N, \quad (5)$$

where the residual is centered and second-order orthogonal to the code nuisance:

$$\mathbb{E}[\eta_N \mid B] = 0, \quad \mathbb{E}[U_C \eta_N^\top \mid B] = 0. \quad (6)$$

The residual nuisance η_N represents additional paraphrase, verbosity, ordering, style, and prompt-form variation in native natural language. Our modelling assumes that once we condition on the core sufficient latent, the residual nuisance from natural language provides no additional information about the code nuisance. This matches prior literature and our intuition. Code has a more constrained grammar and conventional structure, whereas natural-language explanations admit more paraphrastic and discourse-level variation. Work on the naturalness of software likewise treats code as a structured statistical object with regularities distinct from ordinary prose (Allamanis et al., 2018).

Remark 3.2 (Finite-sample OLS learning intuition). For intuition on why the nuisance coordinates hurt learning in the finite sample regime. Let $p = d_B + d_U$ and consider the Gaussian special case

$$S_r \sim \mathcal{N}(0, \Sigma_r), \quad Y = \beta^{*\top} S_r + \varepsilon, \quad \beta^* = (\theta, 0), \quad (7)$$

with $\varepsilon \sim \mathcal{N}(0, \sigma^2)$. Ordinary least squares and ridge regression have finite-sample excess-risk terms controlled by dimension, conditioning, and noise variance. In the isotropic case $\Sigma_r = I_p$, if $\hat{\beta}_r$ is ordinary least squares from $n > p + 1$ independent route- r samples, then

$$\mathbb{E}_{D_{r,n}}[R_r(\hat{\beta}_r)] - R^* = \sigma^2 \frac{p}{n - p - 1}, \quad (8)$$

where $R^* = \sigma^2$. Fitting irrelevant coordinates increases finite-sample prediction error (Hastie et al., 2009; Wainwright, 2019).

3.3.2. RESULTS

Covariance domination. Let

$$\Sigma_{U,r} = \mathbb{E}[U_r U_r^\top] \quad (9)$$

be the nuisance covariance. We use the Loewner order on positive semidefinite matrices: $A \preceq B$ means $B - A$

is positive semidefinite, equivalently $v^\top A v \leq v^\top B v$ for every vector v (Horn & Johnson, 2012).

Definition 3.3 (Nuisance covariance dominance). The code route has no larger nuisance covariance than native natural language if

$$\Sigma_{U,C} \preceq \Sigma_{U,N}. \quad (10)$$

This term is introduced here; the underlying Loewner order is standard.

Proposition 3.4 (Extra nuisance implies covariance dominance). *Under (4) and Assumption 3.1,*

$$\Sigma_{U,C} \preceq \Sigma_{U,N}. \quad (11)$$

Proof. Using (5),

$$\Sigma_{U,N} = \mathbb{E}[(U_C + \eta_N)(U_C + \eta_N)^\top] \quad (12)$$

$$= \mathbb{E}[U_C U_C^\top] + \mathbb{E}[\eta_N \eta_N^\top] + \mathbb{E}[U_C \eta_N^\top] + \mathbb{E}[\eta_N U_C^\top]. \quad (13)$$

The cross terms vanish by the tower property and the second-order orthogonality condition:

$$\mathbb{E}[U_C \eta_N^\top] = \mathbb{E}[\mathbb{E}[U_C \eta_N^\top \mid B]] = 0, \quad (14)$$

and similarly for its transpose. Hence

$$\Sigma_{U,N} - \Sigma_{U,C} = \mathbb{E}[\eta_N \eta_N^\top] \succeq 0. \quad (15)$$

This is the covariance addition formula for conditionally uncorrelated components, written in positive-semidefinite order (Horn & Johnson, 2012; Kallenberg, 2021). $\square$

Uniform linear-risk theorem. The covariance order gives the desired uniform statement for the unrestricted linear class.

Theorem 3.5 (Uniform representation non-inferiority under covariance dominance). *Assume (3), (4), and (10). Also assume $\mathbb{E}[BU_r^\top] = 0$ and $\mathbb{E}[\varepsilon U_r] = 0$ for $r \in \{C, N\}$. Then, for squared loss and the common class $\mathcal{H}_{\text{lin}}$ in (1),*

$$\sup_{h \in \mathcal{H}_{\text{lin}}} \{R_C(h) - R_N(h)\} \leq 0. \quad (16)$$

Consequently code is representation-non-inferior to native natural language with margin $\tau_{\text{rep}} = 0$ in this linear model.

Proof. Fix $h_{a,v} \in \mathcal{H}_{\text{lin}}$. By (3),

$$h_{a,v}(B, U_r) - Y = (a - \theta)^\top B + v^\top U_r - \varepsilon. \quad (17)$$

Using $\mathbb{E}[BU_r^\top] = 0$, $\mathbb{E}[\varepsilon U_r] = 0$, and (4), the terms involving U_r separate:

$$R_r(h_{a,v}) = R_{\text{core}}(a) + v^\top \Sigma_{U,r} v, \quad (18)$$

where $R_{\text{core}}(a) = \mathbb{E}[(a - \theta)^\top B - \varepsilon]^2$ does not depend on r . Therefore

$$R_C(h_{a,v}) - R_N(h_{a,v}) = v^\top (\Sigma_{U,C} - \Sigma_{U,N}) v \leq 0 \quad (19)$$

by (10). Taking the supremum over $h \in \mathcal{H}_{\text{lin}}$ proves (16). $\square$

Corollary 3.5.1 (Approximate covariance dominance on a bounded nuisance class). *Under the assumptions of Theorem 3.5, suppose exact covariance dominance is relaxed to*

$$\Sigma_{U,C} \preceq \Sigma_{U,N} + \kappa I_{d_U} \quad (20)$$

for some $\kappa \geq 0$. Fix a finite nuisance-sensitivity radius $L < \infty$ and restrict the predictor class to

$$\mathcal{H}_L = \{h_{a,v} \in \mathcal{H}_{\text{lin}} : \|v\|_2 \leq L\}. \quad (21)$$

Then

$$\sup_{h \in \mathcal{H}_L} \{R_C(h) - R_N(h)\} \leq L^2 \kappa. \quad (22)$$

Thus code is uniformly representation-non-inferior with margin $\tau_{\text{rep}} = L^2 \kappa$ on the bounded class $\mathcal{H}_L$. Without the bound $\|v\|_2 \leq L$, approximate covariance domination alone does not give a finite uniform margin.

Proof. For $h_{a,v} \in \mathcal{H}_L$, the proof of Theorem 3.5 gives

$$R_C(h_{a,v}) - R_N(h_{a,v}) = v^\top (\Sigma_{U,C} - \Sigma_{U,N}) v. \quad (23)$$

By (20), this is at most $\kappa \|v\|_2^2$, and $\|v\|_2 \leq L$ by definition of $\mathcal{H}_L$. Taking the supremum over $h \in \mathcal{H}_L$ proves (22). $\square$

In the case where the exact version is too strong, the approximate version shows that we can relax the decomposition and the non-inferiority results still hold.

Overall, this shows that representation explains little in the worst-case risk sense.

3.4. Decision Intervention Reduction

When the theory breaks down due to non-linearity and high-dimensionality, we find that the non-inferiority assertion still holds in practice. The implications are that representation does not explain as much of the observed performance uplift as execution.

Intervention setup. Our goal is to provide evidence that code is non-inferior to NL in expected loss in the LLM regime. In our proof, we modelled NL as a fixed transformation (linearly additive projection) of a code latent. Similarly, we apply a fixed transformation of the code (i.e. permutations, syntactic shifts etc) to simulate NL. We expect it to yield nothing worse, since the underlying assumption is that code contains all decision-relevant information. The following experiment demonstrates that this is indeed the case.

We show that a fixed transformation of code (to natural language) behaves functionally similar to original natural language reasoning. Here, our estimand is the final accuracy, and treatment is the different inputs $\in \{\text{translated natural language reasoning, original natural language reasoning}\}$ to our LLM executor.

For a task instance x , we prompt a target language model in three conditions: (1) Baseline: x , (2) Native: $x \parallel z_{\text{NL}}$, (3) Translated: $x \parallel \hat{z}_{\text{NL}}$, where $z_{\text{NL}} \sim p_{\text{NL}}(\cdot \mid x)$ is the native Route 1 trace and $\hat{z}_{\text{NL}}$ is obtained by translating the corresponding code to NL using the same translator procedure. If translated NL loses decision-relevant information relative to native NL, we would expect

$$\text{Acc}(x \parallel \hat{z}_{\text{NL}}) < \text{Acc}(x \parallel z_{\text{NL}}).$$

Protocol and models. We run this test on 1000 held-out instances across multiple translator models (Claude-Haiku-4.5, Gemini-2.5-Flash, Mixtral-8x22B-Instruct). Crucially, in this experiment the translator model is matched to the original code generator (i.e., we translate code produced by the same model family), to avoid confounding functional losses with cross-model stylistic mismatch.

Results. We fail to reject the null hypothesis of equal performance, and observe overlapping 95% confidence intervals between the Native and Translated conditions (Figure 5). This suggests that, under our evaluated settings, translating code into NL does not destroy decision-relevant information for downstream answering, consistent with the claim that NL CoT does not systematically add algorithmic advantage beyond what is in code. Qualitatively, we notice that the translated results are quite similar to the originals, though this seems heavily reliant on the translating prompt.

Interpretation. Our empirical results indicate that Condition 1 holds in the settings we study, implying that the best achievable performance using code traces matches that of natural-language traces up to a small error ε . In particular, we find no evidence that natural-language reasoning introduces novel algorithmic strategies beyond those already captured by code representations on algorithmic tasks. As a result, replacing NL traces with code traces in the generation stage does not constitute a performance bottleneck when optimizing the end-to-end reasoning pipeline. Instead, the remaining limitations are best attributed to the execution mechanism rather than the trace representation (Section 4).

4. Route 2 vs. Route 3 Analysis

The key research question in this section is: **Is execution the bottleneck for language-based reasoning?** To address this question, we also ask two sub-questions:

- 1) When code is *correct*, does external code execution (Route 3) have lower expected loss than LLM-based code execution (Route 2)?

- 2) When code is *incorrect*, is the probability that Route 2 outperforms Route 3 small?

We pinpoint execution as the bottleneck (cf. Section 3.3), highlighting how using LLMs to generate code plans and the code itself is an *optimal* (or near-optimal) policy among the 3 routes (Section 3.2). In the 40-task analysis set, Route 3 exceeds Route 2 by +31.47pp. The recovery mass where Route 2 succeeds while Route 3 fails is 1.61%, whereas the execution-win mass where Route 3 succeeds while Route 2 fails is 33.08%.

4.1. Setup

As defined in Section 2, Route 2 and Route 3 share the same trace generator E_{Code} and differ only in the executor. Let $\rho_{\text{Sim}} \in \mathcal{H}_C$ denote the fixed LLM-based simulation executor used in Route 2. Let g be a deterministic interpreter mapping (x, z_C) to $\mathcal{Y} \cup \{\perp\}$, where $\perp$ denotes execution failure. Define the instance-correct execution event

$$C := \{g(X, Z_C) = Y^*(X)\}.$$

Risk decomposition. Let

$$\begin{aligned} e_C &:= \Pr(\hat{Y}_{\text{Sim}} \neq Y^*(X) \mid C), \\ r &:= \Pr(\hat{Y}_{\text{Sim}} = Y^*(X) \mid \neg C). \end{aligned}$$

Then

$$\begin{aligned} R(E_{\text{Code}}, \rho_{\text{Exec}}) &= \Pr(\neg C), \\ R(E_{\text{Code}}, \rho_{\text{Sim}}) &= \Pr(C) e_C + \Pr(\neg C)(1 - r), \end{aligned}$$

and hence

$$R(E_{\text{Code}}, \rho_{\text{Sim}}) - R(E_{\text{Code}}, \rho_{\text{Exec}}) = \Pr(C) e_C - \Pr(\neg C) r.$$

Implications. Simulation noise on correct-execution instances ($\Pr(C) e_C$) harms Route 2, while recovery on incorrect or failed executions ($\Pr(\neg C) r$) helps Route 2. Route 3 dominates whenever the recovery mass is insufficient to offset simulation noise, a condition quantified empirically in Section 3.2.

Empirical recovery as an upper bound. A direct way to operationalize the “recovery” term in the decomposition is to measure the frequency with which Route 2 succeeds while Route 3 fails on the *same* generated code trace, i.e.,

$$\Pr(\hat{Y}_{\text{Sim}} = Y^*(X), g(X, Z_C) \neq Y^*(X)).$$

This quantity corresponds to the *recovery mass* $\Pr(\neg C) r$ appearing in the above risk gap. It aggregates both (i) genuine recovery from incorrect code and (ii) cases where execution fails (e.g., $g(X, Z_C) = \perp$) but the simulator still answers correctly. Because the simulator may partially ignore Z_C and answer directly from X , this statistic should be interpreted as an *upper bound* on mechanistic “recovery from flawed code.”




Figure 6. **Recovery mass (blue) is 1.61% overall, while execution-win mass (red) is 33.08% overall.** Recovery mass $\Pr(-C)r$ counts cases where code execution fails but LLM simulation succeeds. Execution-win mass $\Pr(C)e_C$ counts cases where code execution succeeds and LLM simulation fails.

Interpretation. Route 2 can outperform Route 3 only if the recovery mass $\Pr(-C)r$ is large enough to offset simulation noise on correct-execution instances $\Pr(C)e_C$. Empirically, we find that recovery mass is consistently small across tasks and models (Figure 6), indicating that Route 2 rarely compensates for execution failures via recovery. In the 40-task analysis set, this recovery mass is 1.61%, while the execution-win mass $\Pr(C)e_C$ is 33.08%. This explains why deterministic execution typically achieves lower end-to-end error than simulation on the same generated code traces (Figure 3).

5. Related Work and Discussion

Neuro-symbolic Learning. This paper builds on research in neuro-symbolic integration (Graves et al., 2014; Veličković & Blundell, 2021; Reed & Freitas, 2016; Graves et al., 2016), which combines neural networks with symbolic reasoning systems. These approaches are motivated by cognitive science (Schneider & Chein, 2003; Risko & Gilbert, 2016; Anderson, 2010), hierarchical reinforcement learning (Kolter et al., 2007; Dietterich, 2000), and compositionality research (Hudson & Manning, 2018; Hupkes et al., 2020; Andreas et al., 2017; Poggio et al., 2017). An orthogonal line of work explores direct execution of algorithms by neural networks (Veličković & Blundell, 2021; Mahdavi et al., 2023; Ibarz et al., 2022; Yan et al., 2020). Unlike these approaches that focus on *how* to integrate neural and symbolic components, our work addresses *why* neuro-symbolic integration outperforms neural reasoning alone for algorithmic tasks (Sections 3.3 and 4).

LLM Reasoning. Recent work has explored various reasoning paradigms for LLMs, including symbolic reasoning (Marra et al., 2019; Olausson et al., 2023; Han et al., 2024), chain-of-thought prompting (Wei et al., 2022; Zelikman et al., 2022; Merrill & Sabharwal, 2024; Altabaa et al., 2025), and in-context learning (Xie et al., 2021; Garg

et al., 2022; Akyürek et al., 2022; Zhang et al., 2024). Xie et al. (2021) model in-context learning as implicit Bayesian inference, which we extend to compare different reasoning representations. While prior work demonstrates *that* certain prompting strategies improve performance, we provide a theoretical framework (Section 2) explaining *why* code representations lead to lower Bayes error in our representation analysis (Section 3.3).

LLM Tool-Use. Tool-augmented LLMs have achieved strong empirical results (Shen, 2024; Schick et al., 2023; Qin et al., 2023; Tang et al., 2023; Parisi et al., 2022). Code generation for tool-use can be viewed as a form of semantic parsing (Shin & Durme, 2022; Krishnamurthy et al., 2017; Berant et al., 2013; Dong & Lapata, 2016) or function calling (Puri et al., 2021; Alon et al., 2019; Chen & Zhou, 2018). Our work complements this literature by providing theoretical justification (Sections 3.3 and 4) for the observed empirical advantages of code-based tool-use over direct natural language reasoning.

6. Conclusion

We introduced a three-route framework (Section 2) for comparing code and natural language representations in algorithmic reasoning. By introducing an intermediate route—code generation with LLM-based simulation—we isolate translation effects from execution effects, enabling tractable pairwise analysis. Empirically, code execution achieves +31.6pp v.s. NL across the 40-task analysis set and 6 models (Section 3.2), while code simulation and NL reasoning remain statistically close. In our analysis (Section 3.3), we first show that under the optimal risk regime, code is at least as good as NL up to a small ϵ bound. Thus, in the evaluated setting, trace generation with code is not the observed bottleneck relative to natural language. We further verify in the second part of our analysis that execution is the bottleneck: recovery mass where simulation succeeds but execution fails is only 1.61%, while execution-win mass where execution succeeds but simulation fails is 33.08% (Figure 6). Overall, when designing stronger AI systems that solve harder tasks, understanding that language reasoning struggles to scale with hardness should motivate us to design compositional AI systems that plan and reason in code and delegate execution to external solvers, as opposed to monolithic architectures.

Limitations. (1) *Empirical approximations:* Our distributional similarity tests use finite samples and specific judge models, introducing estimation variance. Furthermore, there is room for optimization with the prompts of both judges and translators. (2) *Model coverage:* While we evaluate both frontier (GPT-4o, Claude, Gemini) and open-source models (Mistral, Codestral), we cannot guarantee findings generalize to all architectures or future models. (3) *Theoretical assumptions:* The statistical analysis result assumes the existence of a garbling channel from code to NL, which we

validate empirically but do not prove from first principles.

7. Impact Statement

This paper presents work whose goal is to advance the field of machine learning. There are many potential societal consequences of our work, none of which we feel must be specifically highlighted here.

References

- Akyürek, E., Schuurmans, D., Andreas, J., Ma, T., and Zhou, D. What learning algorithm is in-context learning? investigations with linear models. *arXiv preprint arXiv:2211.15661*, 2022. URL <https://arxiv.org/abs/2211.15661>.
- Allamanis, M., Barr, E. T., Devanbu, P., and Sutton, C. A survey of machine learning for big code and naturalness. *ACM Computing Surveys*, 51(4):1–37, 2018. doi: 10.1145/3212695.
- Alon, U., Zilberstein, M., Levy, O., and Yahav, E. code2vec: learning distributed representations of code. *Proceedings of the ACM on Programming Languages*, 3(POPL):1–29, January 2019. ISSN 2475-1421. doi: 10.1145/3290353.
- Altabaa, A., Montasser, O., and Lafferty, J. CoT Information: Improved Sample Complexity under Chain-of-Thought Supervision, May 2025. URL <http://arxiv.org/abs/2505.15927>. arXiv:2505.15927 [stat].
- Anderson, M. L. Neural reuse: A fundamental organizational principle of the brain. *Behavioral and Brain Sciences*, 33(4):245–266, August 2010. ISSN 0140-525X, 1469-1825. doi: 10.1017/S0140525X10000853. URL https://www.cambridge.org/core/product/identifier/S0140525X10000853/type/journal_article.
- Andreas, J., Rohrbach, M., Darrell, T., and Klein, D. Neural Module Networks, July 2017. URL <http://arxiv.org/abs/1511.02799>. arXiv:1511.02799 [cs].
- Berant, J., Chou, A., Frostig, R., and Liang, P. Semantic Parsing on Freebase from Question-Answer Pairs. In Yarowsky, D., Baldwin, T., Korhonen, A., Livescu, K., and Bethard, S. (eds.), *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pp. 1533–1544, Seattle, Washington, USA, October 2013. Association for Computational Linguistics. URL <https://aclanthology.org/D13-1160/>.
- Chen, Q. and Zhou, M. A neural framework for retrieval and summarization of source code. In *Proceedings of the 33rd ACM/IEEE International Conference on Automated Software Engineering*, pp. 826–831, Montpellier France, September 2018. ACM. ISBN 978-1-4503-5937-5. doi: 10.1145/3238147.3240471. URL <https://dl.acm.org/doi/10.1145/3238147.3240471>.
- Dietterich, T. G. Hierarchical Reinforcement Learning with the MAXQ Value Function Decomposition. *Journal of Artificial Intelligence Research*, 13: 227–303, November 2000. ISSN 1076-9757. doi: 10.1613/jair.639. URL <https://www.jair.org/index.php/jair/article/view/10266>.
- Dong, L. and Lapata, M. Language to Logical Form with Neural Attention, June 2016. URL <http://arxiv.org/abs/1601.01280>. arXiv:1601.01280 [cs].
- Gao, L., Madaan, A., Zhou, S., Alon, U., Liu, P., Yang, Y., Callan, J., and Neubig, G. PAL: Program-aided language models. In *Proceedings of the 40th International Conference on Machine Learning*, pp. 10764–10799. PMLR, 2023. URL <https://proceedings.mlr.press/v202/gao23f.html>.
- Garg, S., Tsipras, D., Liang, P. S., and Valiant, G. What can transformers learn in-context? a case study of simple function classes. *Advances in neural information processing systems*, 35:30583–30598, 2022. URL <https://arxiv.org/abs/2208.01066>.
- Geirhos, R., Jacobsen, J.-H., Michaelis, C., Zemel, R., Brendel, W., Bethge, M., and Wichmann, F. A. Short-cut learning in deep neural networks. *Nature Machine Intelligence*, 2:665–673, 2020. doi: 10.1038/s42256-020-00257-z.
- Graves, A., Wayne, G., and Danihelka, I. Neural Turing Machines, December 2014. URL <http://arxiv.org/abs/1410.5401>. arXiv:1410.5401 [cs].
- Graves, A., Wayne, G., Reynolds, M., Harley, T., Danihelka, I., Grabska-Barwińska, A., Colmenarejo, S. G., Grefenstette, E., Ramalho, T., Agapiou, J., Badia, A. P., Hermann, K. M., Zwols, Y., Ostrovski, G., Cain, A., King, H., Summerfield, C., Blunsom, P., Kavukcuoglu, K., and Hassabis, D. Hybrid computing using a neural network with dynamic external memory. *Nature*, 538 (7626):471–476, October 2016. ISSN 0028-0836, 1476-4687. doi: 10.1038/nature20101. URL <https://www.nature.com/articles/nature20101>.
- Gururangan, S., Swayamdipta, S., Levy, O., Schwartz, R., Bowman, S. R., and Smith, N. A. Annotation artifacts in natural language inference data. In *Proceedings of NAACL-HLT*, 2018. URL <https://aclanthology.org/N18-2017/>.

- Han, S., Schoelkopf, H., Zhao, Y., Qi, Z., Riddell, M., Zhou, W., Coady, J., Peng, D., Qiao, Y., Benson, L., Sun, L., Wardle-Solano, A., Szabo, H., Zubova, E., Burtell, M., Fan, J., Liu, Y., Wong, B., Sailor, M., Ni, A., Nan, L., Kasai, J., Yu, T., Zhang, R., Fabbri, A. R., Kryscinski, W., Yavuz, S., Liu, Y., Lin, X. V., Joty, S., Zhou, Y., Xiong, C., Ying, R., Cohan, A., and Radev, D. FOLIO: Natural Language Reasoning with First-Order Logic, October 2024. URL <http://arxiv.org/abs/2209.00840>. arXiv:2209.00840 [cs].
- Hastie, T., Tibshirani, R., and Friedman, J. *The Elements of Statistical Learning*. Springer, 2 edition, 2009. URL <https://hastie.su.domains/ElemStatLearn/>.
- Horn, R. A. and Johnson, C. R. *Matrix Analysis*. Cambridge University Press, 2 edition, 2012. doi: 10.1017/CBO9781139020411.
- Hudson, D. A. and Manning, C. D. Compositional Attention Networks for Machine Reasoning, April 2018. URL <http://arxiv.org/abs/1803.03067>. arXiv:1803.03067 [cs].
- Hupkes, D., Dankers, V., Mul, M., and Bruni, E. Compositionality decomposed: how do neural networks generalise?, February 2020. URL <http://arxiv.org/abs/1908.08351>. arXiv:1908.08351 [cs].
- Ibarz, B., Kurin, V., Papamakarios, G., Nikiforou, K., Benani, M., Csordás, R., Dudzik, A. J., Bošnjak, M., Vitvitskyi, A., Rubanova, Y., Deac, A., Bevilacqua, B., Ganin, Y., Blundell, C., and Veličković, P. A Generalist Neural Algorithmic Learner. In *Proceedings of the First Learning on Graphs Conference*, pp. 2:1–2:23. PMLR, December 2022. URL <https://proceedings.mlr.press/v198/ibarz22a.html>. ISSN: 2640-3498.
- Kallenberg, O. *Foundations of Modern Probability*. Springer, 3 edition, 2021. doi: 10.1007/978-3-030-61871-1.
- Kolter, J., Abbeel, P., and Ng, A. Hierarchical Apprenticeship Learning with Application to Quadruped Locomotion. In *Advances in Neural Information Processing Systems*, volume 20. Curran Associates, Inc., 2007. URL https://proceedings.neurips.cc/paper_files/paper/2007/hash/54a367d629152b720749e187b3eaa11b-Abstract.html.
- Krishnamurthy, J., Dasigi, P., and Gardner, M. Neural Semantic Parsing with Type Constraints for Semi-Structured Tables. In Palmer, M., Hwa, R., and Riedel, S. (eds.), *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing*, pp. 1516–1526, Copenhagen, Denmark, September 2017. Association for Computational Linguistics. doi: 10.18653/v1/D17-1160. URL <https://aclanthology.org/D17-1160/>.
- Lehmann, E. L. and Casella, G. *Theory of Point Estimation*. Springer, 2 edition, 1998. doi: 10.1007/b98854.
- Lu, Y., Bartolo, M., Moore, A., Riedel, S., and Stenetorp, P. Fantastically ordered prompts and where to find them: Overcoming few-shot prompt order sensitivity. In *Proceedings of ACL*, 2022. URL <https://aclanthology.org/2022.acl-long.556/>.
- Lyu, Q., Havaladar, S., Stein, A., Zhang, L., Rao, D., Wong, E., Apidianaki, M., and Callison-Burch, C. Faithful chain-of-thought reasoning. In *The 13th International Joint Conference on Natural Language Processing and the 3rd Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics (IJCNLP-AACL 2023)*, pp. 305–329. Association for Computational Linguistics, 2023. doi: 10.18653/v1/2023.ijcnlp-main.20.
- Mahdavi, S., Swersky, K., Kipf, T., Hashemi, M., Thrampoulidis, C., and Liao, R. Towards Better Out-of-Distribution Generalization of Neural Algorithmic Reasoning Tasks, March 2023. URL <http://arxiv.org/abs/2211.00692>. arXiv:2211.00692 [cs].
- Marra, G., Giannini, F., Diligenti, M., and Gori, M. Integrating Learning and Reasoning with Deep Logic Models, January 2019. URL <http://arxiv.org/abs/1901.04195>. arXiv:1901.04195 [cs].
- McCoy, R. T., Pavlick, E., and Linzen, T. Right for the wrong reasons: Diagnosing syntactic heuristics in natural language inference. In *Proceedings of ACL*, 2019. URL <https://aclanthology.org/P19-1334/>.
- Merrill, W. and Sabharwal, A. The Expressive Power of Transformers with Chain of Thought, April 2024. URL <http://arxiv.org/abs/2310.07923>. arXiv:2310.07923 [cs].
- Olausson, T. X., Gu, A., Lipkin, B., Zhang, C. E., Solar-Lezama, A., Tenenbaum, J. B., and Levy, R. LINC: A Neurosymbolic Approach for Logical Reasoning by Combining Language Models with First-Order Logic Provers. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pp. 5153–5176, 2023. doi: 10.18653/v1/2023.emnlp-main.313. URL <http://arxiv.org/abs/2310.15164>. arXiv:2310.15164 [cs].
- Pan, L., Albalak, A., Wang, X., and Wang, W. Y. LogicIm: Empowering large language models with symbolic

- solvers for faithful logical reasoning. *arXiv preprint arXiv:2305.12295*, 2023. URL <https://arxiv.org/abs/2305.12295>.
- Parisi, A., Zhao, Y., and Fiedel, N. TALM: Tool Augmented Language Models, May 2022. URL <http://arxiv.org/abs/2205.12255>. arXiv:2205.12255 [cs].
- Pezeshkpour, P. and Hruschka, E. Large language models sensitivity to the order of options in multiple-choice questions. *arXiv preprint arXiv:2308.11483*, 2023. URL <https://arxiv.org/abs/2308.11483>.
- Poggio, T., Mhaskar, H., Rosasco, L., Miranda, B., and Liao, Q. Why and when can deep-but not shallow-networks avoid the curse of dimensionality: a review. *International Journal of Automation and Computing*, 14(5):503–519, 2017. doi: 10.1007/s11633-017-1054-2.
- Puri, R., Kung, D. S., Janssen, G., Zhang, W., Domeniconi, G., Zolotov, V., Dolby, J., Chen, J., Choudhury, M., Decker, L., Thost, V., Buratti, L., Pujar, S., Ramji, S., Finkler, U., Malaika, S., and Reiss, F. CodeNet: A Large-Scale AI for Code Dataset for Learning a Diversity of Coding Tasks, August 2021. URL <http://arxiv.org/abs/2105.12655>. arXiv:2105.12655 [cs].
- Qin, Y., Liang, S., Ye, Y., Zhu, K., Yan, L., Lu, Y., Lin, Y., Cong, X., Tang, X., Qian, B., Zhao, S., Hong, L., Tian, R., Xie, R., Zhou, J., Gerstein, M., Li, D., Liu, Z., and Sun, M. ToolLLM: Facilitating Large Language Models to Master 16000+ Real-world APIs, October 2023. URL <http://arxiv.org/abs/2307.16789>. arXiv:2307.16789 [cs].
- Reed, S. and Freitas, N. d. Neural Programmer-Interpreters, February 2016. URL <http://arxiv.org/abs/1511.06279>. arXiv:1511.06279 [cs].
- Risko, E. F. and Gilbert, S. J. Cognitive Offloading. *Trends in Cognitive Sciences*, 20(9):676–688, September 2016. ISSN 13646613. doi: 10.1016/j.tics.2016.07.002. URL <https://linkinghub.elsevier.com/retrieve/pii/S1364661316300985>.
- Schick, T., Dwivedi-Yu, J., Dessí, R., Raileanu, R., Lomeli, M., Hambro, E., Zettlemoyer, L., Cancedda, N., and Scialom, T. Toolformer: language models can teach themselves to use tools. In *Proceedings of the 37th International Conference on Neural Information Processing Systems*, NIPS ’23, Red Hook, NY, USA, 2023. Curran Associates Inc. URL <https://arxiv.org/abs/2302.04761>.
- Schneider, W. and Chein, J. M. Controlled & automatic processing: behavior, theory, and biological mechanisms. *Cognitive Science*, 27(3):525–559, May 2003. ISSN 0364-0213, 1551-6709. doi: 10.1207/s15516709cog2703_8.
- Shen, Z. LLM With Tools: A Survey, September 2024. URL <http://arxiv.org/abs/2409.18807>. arXiv:2409.18807 [cs].
- Shin, R. and Durme, B. V. Few-Shot Semantic Parsing with Language Models Trained On Code, May 2022. URL <http://arxiv.org/abs/2112.08696>. arXiv:2112.08696 [cs].
- Tang, Q., Deng, Z., Lin, H., Han, X., Liang, Q., Cao, B., and Sun, L. ToolAlpaca: Generalized Tool Learning for Language Models with 3000 Simulated Cases, September 2023. URL <http://arxiv.org/abs/2306.05301>. arXiv:2306.05301 [cs].
- Veličković, P. and Blundell, C. Neural algorithmic reasoning. *Patterns*, 2(7):100273, July 2021. ISSN 26663899. doi: 10.1016/j.patter.2021.100273. URL <https://linkinghub.elsevier.com/retrieve/pii/S2666389921000994>.
- Wainwright, M. J. *High-Dimensional Statistics: A Non-Asymptotic Viewpoint*. Cambridge University Press, 2019. doi: 10.1017/9781108627771.
- Wang, Z., Chang, Q., Patel, H., Biju, S., Wu, C.-E., Liu, Q., Ding, A., Rezazadeh, A., Shah, A., Bao, Y., and Siow, E. Mcp-bench: Benchmarking tool-using llm agents with complex real-world tasks via mcp servers, 2025. URL <https://arxiv.org/abs/2508.20453>.
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Xia, F., Chi, E., Le, Q. V., Zhou, D., et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35: 24824–24837, 2022. URL <https://arxiv.org/abs/2201.11903>.
- Xie, S. M., Raghunathan, A., Liang, P., and Ma, T. An explanation of in-context learning as implicit bayesian inference. *arXiv preprint arXiv:2111.02080*, 2021. URL <https://arxiv.org/abs/2111.02080>.
- Yan, Y., Swersky, K., Koutra, D., Ranganathan, P., and Hashemi, M. Neural Execution Engines: Learning to Execute Subroutines, October 2020. URL <http://arxiv.org/abs/2006.08084>. arXiv:2006.08084 [cs].
- Yang, J., Jimenez, C. E., Wettig, A., Lieret, K., Yao, S., Narasimhan, K. R., and Press, O. SWE-agent: Agent-computer interfaces enable automated software engineering. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024. URL <https://openreview.net/forum?id=mXpq6ut8J3>.

- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., and Cao, Y. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023. URL https://openreview.net/forum?id=WE_vluYUL-X.
- Zelikman, E., Wu, Y., Mu, J., and Goodman, N. D. STaR: Bootstrapping Reasoning With Reasoning, May 2022. URL <http://arxiv.org/abs/2203.14465>. arXiv:2203.14465 [cs].
- Zhang, R., Frei, S., and Bartlett, P. L. Trained transformers learn linear models in-context. *Journal of Machine Learning Research*, 25(49):1–55, 2024. URL <https://jmlr.org/papers/v25/23-1042.html>.
- Zhao, T. Z., Wallace, E., Feng, S., Klein, D., and Singh, S. Calibrate before use: Improving few-shot performance of language models. In *Proceedings of the International Conference on Machine Learning*, 2021. URL <https://proceedings.mlr.press/v139/zhao21c.html>.

A. Supplementary Route Results

Route definitions. Across all route tables, Route 1 denotes natural-language reasoning, Route 2 denotes natural-language simulation over a code trace, and Route 3 denotes deterministic code execution. Captions and result notes state when a table uses the full evaluation set, a 25% subset, a single seed, or per-arm parse-normalized denominators. Unless otherwise stated, accuracies are percentages and route differences are percentage points.

A.1. Complexity-Stratified Route Accuracy

We stratify the main route-accuracy evaluation by the asymptotic complexity class of each task. Across classes, Route 3 remains the main source of improvement over Route 2, while Route 2 stays close to Route 1.

Complexity	Tasks	Inst.	R1	R2	R3	R2-R1	$p_{2,1}$	R3-R2
$O(1)$	3	180	46.2	47.3	86.2	1.0	0.0206	39.0
$O(\log n)$	1	35	22.4	28.6	32.9	6.2	2.7×10^{-4}	4.3
$O(n)$	4	135	4.8	5.0	7.5	0.2	0.635	2.5
$O(n \log n)$	5	68	0.0	0.0	0.0	0.0	1	0.0
$O(n^2)$	16	297	10.3	10.5	38.9	0.1	0.644	28.5
$O(n^2 \log n)$	1	1	0.0	0.0	0.0	0.0	1	0.0
$O(n^3)$	4	37	0.0	0.0	0.0	0.0	1	0.0
NP-hard	6	360	17.6	16.8	69.7	-0.8	0.0278	53.0

Table 1. Route accuracy grouped by asymptotic complexity in the main route evaluation. Tasks and Inst. report task and instance counts; $p_{2,1}$ is the exact two-sided McNemar test for Route 2 vs. Route 1 on paired outcomes.

Result. Table 1 shows the aggregate route pattern within the higher-support complexity strata: Route 3–Route 2 is the largest observed route change, while Route 2 stays close to Route 1.

Task mapping. The 40-task analysis set excludes `segments_intersect`, `knap`, `gcp`, and `spp` from the broader 44-task route-analysis mapping. The largest retained strata are $O(n^2)$ dynamic-programming, sorting, string, graph, and shortest-path tasks; and NP-hard ILP, edge-disjoint paths, shortest-path, and TSP tasks. The smaller strata cover constant-time arithmetic, logarithmic binary search, linear selection/string matching, $O(n \log n)$ scheduling/sorting/hulls, $O(n^2 \log n)$ Kruskal MST, and cubic dynamic-programming/shortest-path routines.

A.2. Model-Stratified Route Accuracy

We stratify the main route comparison by model to check whether the aggregate result is driven by one model family or scale regime. The Route 3 advantage persists across all six full-coverage models, while the Route 2–Route 1 difference is small and changes sign across models.

Model	Type	Inst.	R1	R2	R3	R2-R1	$p_{2,1}$	R3-R2	$p_{3,2}$
anthropic/claude-haiku-4.5	closed	1113	23.3	23.6	47.1	0.3	0.612	23.5	2.4×10^{-148}
google/gemini-2.0-flash-001	closed	1113	17.5	18.1	51.8	0.7	0.137	33.7	8.5×10^{-298}
google/gemini-2.5-flash	closed	1113	20.7	20.5	50.2	-0.1	0.823	29.7	3.1×10^{-238}
openai/gpt-4o-mini	closed	1113	14.1	12.9	49.1	-1.2	0.00663	36.1	$< 10^{-300}$
mistralai/codestral-2508	open	1113	15.0	15.8	52.7	0.8	0.0279	36.9	$< 10^{-300}$
mistralai/mixtral-8x22b-instruct	open	1113	12.7	13.2	42.1	0.5	0.127	28.9	7.2×10^{-253}

Table 2. Route accuracy grouped by model on the 40-task analysis set. Each row contains 1,113 unique problems and 3,339 paired route evaluations; $p_{2,1}$ and $p_{3,2}$ are exact two-sided McNemar tests for Route 2 vs. Route 1 and Route 3 vs. Route 2, respectively.

Result. Table 2 shows that Route 3 stays above Route 2 for every full-coverage model row, so the aggregate execution gain is not driven by a single weak or strong model.

B. Functional Similarity Checks

B.1. Prompt Translation Shot Ablation

The functional similarity test depends on translating code traces back into natural language. To test sensitivity to prompting, we vary the number of in-context examples used by the translator. More examples improve translated-trace utility and narrow the gap to native natural-language traces.

Is Code Better Than Language for Algorithmic Reasoning?

Shots	x	x + native NL	x + translated NL	Δ native	Δ translated	Gap
0	32.70%	55.70%	42.41%	+23.00pp	+9.70pp	-13.29pp
1	32.70%	55.70%	45.78%	+23.00pp	+13.08pp	-9.92pp
2	32.70%	55.70%	49.58%	+23.00pp	+16.88pp	-6.12pp
3	32.70%	55.70%	48.52%	+23.00pp	+15.82pp	-7.17pp
4	32.70%	55.70%	49.37%	+23.00pp	+16.67pp	-6.33pp
5	32.70%	55.70%	50.00%	+23.00pp	+17.30pp	-5.70pp
10 (reference)	39.00%	56.50%	52.00%	+17.50pp	+13.00pp	-4.50pp

Table 3. Translation-additivity shot ablation for Claude Haiku 4.5. Gap is translated NL minus native NL. Rows 0–5 use a matched 25% subset sweep; the 10-shot row is an original main-evaluation reference and is not from the same subset sweep.

Result. Table 3 shows that translated traces improve over the question-only baseline at every reported shot setting, but remain below native natural-language traces in every row.

C. Additional Model Evaluations

C.1. Recursive Language Model Evaluation

To broaden coverage beyond standard LLM forward passes, we evaluate recursive language model execution on a 25% subset at seed 0. The code and natural-language arms remain close under this executor family.

Model	RLM Code Acc	RLM NL Acc	Better Arm
Claude Haiku 4.5	29.4%	30.9%	NL
Gemini 2.5 Pro	47.5%	45.2%	Code

Table 4. Recursive language model code-vs.-NL accuracy on the 25% subset, seed 0. Better Arm marks the higher row accuracy.

Result. Table 4 has no uniform winner: the NL arm is higher for Claude Haiku 4.5, while the code arm is higher for Gemini 2.5 Pro. The corresponding counts are 202/686 code-correct vs. 212/686 NL-correct for Claude Haiku 4.5, and 326/686 code-correct vs. 310/686 NL-correct for Gemini 2.5 Pro. Because this is a 25% subset, we use it as model-coverage evidence rather than as the main route estimate.

C.2. Coding-Specialized Model Evaluation

We evaluate coding-specialized models to test whether models trained on code show a different Route 2 pattern. Even for these models, replacing natural-language traces with code traces changes little compared with replacing LLM simulation by actual code execution.

Model	NL	Sim	Code Exec
x-ai/grok-code-fast-1 (25% data)	47.71% (167/350)	47.71% (167/350)	55.99% (159/284)
qwen/qwen3-coder (25% data)	30.23% (104/344)	26.86% (94/350)	56.19% (168/299)
codestral-2508 (original)	19.89% (943/4740)	23.14% (1097/4740)	59.65% (2266/3799)

Table 5. Route accuracy for coding-specialized models. Each cell reports accuracy with correct/denominator counts; denominators are parse-normalized per arm.

Result. Table 5 shows Code Exec as the highest-accuracy arm in each coding-specialized row. Grok and Qwen use 25% subset runs; Codestral uses the original evaluation run. Rows use per-arm parse-normalized denominators and are not pooled into the main six-model estimate.

C.3. Frontier Model Controlled-Subset Evaluation

We also evaluate two frontier models on a 350-instance subset at seed 1 with tool calling disabled and no patching or repair step. These results are not used as the main estimate because they are smaller than the full six-model evaluation, but they provide a controlled check that the route pattern is not explained only by explicit tool calls.

Model	Route 1	Route 2	Route 3
GPT-5.4	42.57% (149/350)	41.43% (145/350)	54.01% (175/324)
Claude Opus 4.6	52.73% (174/330)	73.42% (116/158)	77.54% (107/138)

Table 6. Frontier tool-disabled controlled-subset route accuracy on the seed 1 350-instance subset with no patching or repair. Entries report accuracy with correct/denominator counts.

Result. Table 6 keeps Route 3 ahead in both tool-disabled subset rows. Reasoning-mode runs are excluded because hidden reasoning and tool scaffolding would confound the route comparison.

D. Failure Analysis

D.1. Code Execution Failure Modes

Among Route 3 code-execution failures, most failures are semantic wrong answers rather than syntax or timeout failures. This supports the interpretation that these errors often originate in the generated program specification rather than in the Python runtime itself.

Model	Wrong answer	Syntax error	Runtime error	Time limit
Haiku	56.26%	10.07%	1.88%	31.80%
Codestral	81.80%	5.87%	9.00%	3.33%
Gemini 2.0	77.22%	14.65%	1.86%	6.27%
Gemini 2.5	77.63%	17.10%	1.82%	3.45%
GPT-4o mini	72.93%	4.56%	18.97%	3.54%
Mixtral	60.67%	4.95%	32.14%	2.24%
Total	70.43%	9.34%	11.55%	8.67%

Table 7. Failure-mode distribution conditional on Route 3 code-execution failures. Rows are within-model percentages and sum to 100% up to rounding.

Result. Table 7 is conditional on Route 3 code-execution failures and excludes parse-only rows where execution completed but the returned value failed parsing or type checks. This keeps the diagnostic focused on what happened after code execution was actually attempted.